


# A Performance Evaluation of Open Source Graph Databases

**Robert McColl**  
**Jason Poovey**

**David Ediger**  
**Dan Campbell**

**David A. Bader**




**Georgia  
Tech**




College of  
Computing


Computational Science and Engineering



# Overview

- Motivation
- Options
- Evaluation
- Results
- Lessons Learned
- Moving Forward





# Massive Streaming Semantic Graphs

## Features

- Millions to **billions** of vertices and edges with rich semantic information (name, type, weight, time), possibly missing or inconsistent data
- Thousands to **millions** of updates per second
- Power-law degree distribution, sparse ( $d(v) \sim O(1)$ ), low diameter



## Financial

- NYSE processes **>2TB daily**, maintains 10's of PB



## Social

- **50,000+** Facebook Likes per second, **1.2B** users
- **6,000** Twitter Tweets per second, **500M** users



## Google

- “Several dozen” **1PB** data sets
- Knowledge Graph: **570M** entities, **18B** relationships



## Business

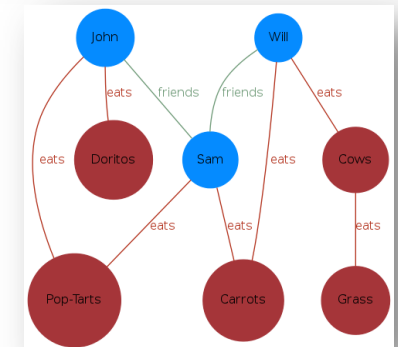
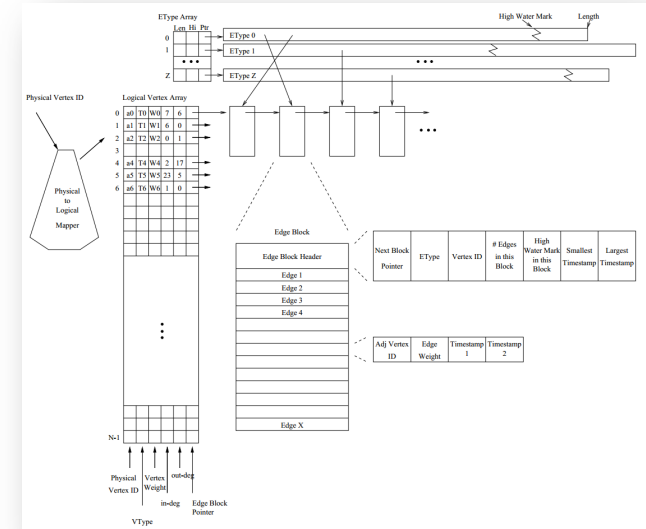
- eBay: **>17** trillion records, **5B** new records per day






# Given these graphs...

- How do we **store** them?
  - Memory vs. Disk
  - ACID vs. loose consistency
  - Single node vs. distributed
  - Semantic and temporal information
  - Numeric IDs
- How do we **query** them?
  - Simple neighbor queries
  - Traversal methods, filtering
  - Programming models vs. query languages
- How do we **compute** meaningful answers?
  - Algorithms and implementation paradigms
  - BSP, MapReduce, MPI, OpenMP, single threaded
  - C, C++, Python, Java, Scala, Javascript, SQL, Gremlin






# The Options

	Owner / Maintainer	License	Platform	Language	Distribution	Cost	Transactional	Memory-Based	Disk-Based	Single-Node	Distributed	Graph Algorithms	Text-Based Query Language	Embeddable	Software	Data Store	Type	
MySQL	Oracle	GPL/Proprietary	x86	C/C++	Bin	Free	X	-	X	X	-	-	SQL	X	X	X	SQLDB	
Oracle	Oracle	Proprietary	x86	C/C++	Bin	\$180-\$950	X	-	X	X	X	-	SQL	X	X	X	SQLDB	
SQL Server	Microsoft	Proprietary	x86-Win	C++	Bin	\$898-\$8592	X	-	X	X	-	-	SQL	-	X	X	SQLDB	
SQLite	Richard Hipp	Public Domain	x86	C	Src/Bin	Free	X	X	X	X	-	-	SQL	X	X	X	SQLDB	
AllegroGraph	Franz, Inc.	Proprietary	x86	Likely Java	Bin	Free-ish/\$\$\$\$	X	-	X	X	-	X	SPARQL,RDFS++,Prolog	-	X	X	GDB	
ArangoDB	ArangoDB	Apache	x86	C/C++/JS	Src/Bin	Free	-	-	X	X	-	-	AQL	-	X	X	GDB/KV/DOC	
DEX	Sparsity-Technologies	Proprietary	x86	C++	Bin	Free Personal/Commercial \$\$	X	-	X	X	-	X	Traversal	X	-	X	GDB	
FlockDB	Twitter	Apache	x86	Java,Scala,Ruby	Src	Free	-	-	X	X	X	-	-	-	-	X	X	GDB
GraphBase	FactNexus	Proprietary	x86	Java	Bin	Free,\$15/mo,\$20,000	?	-	X	X	-	-	Bounds	X	X	X	GDB	
HyperGraphDB	Kobrix Software	LGPL	x86	Java	Src	Free	MVCC	X	X	X	X	-	HGQuery,Traversal	X	-	-	HyperGDB	
InfiniteGraph	Objectivity	Proprietary	x86	Java/C++	Bin	Free Trial/\$5,000	Both	-	X	X	X	-	Gremlin	X	X	X	GDB	
InfoGrid	Johannes Ernst	AGPL/Proprietary	x86	Java	Src/Bin	Free + Support	-	X	X	X	X	-	-	X	X	-	GDB	
Neo4j	Neo Technology	GPL/Proprietary	x86	Java	Src/Bin	Free,\$6,000-\$24,000	X	-	X	X	-	X	Cypher	X	X	X	GDB/NoSQL	
OrientDB	NuvolaBase Ltd	Apache	x86	Java	Src/Bin	Free + Support	Both	X	X	X	X	-	Extended SQL, Gremlin	X	X	X	GDB/NoSQL	
Titan	Aurelius	Apache	x86	Java	Src/Bin	Free + Support	Both	-	X	X	X	-	Gremlin	X	X	-	GDB	
Bagel	UC Berkley	BSD	x86	Java/Scala/Spark	Src	Free	-	X	-	X	X	X	-	X	-	-	BSP	
BGL	Boost / IU	Boost	x86 / C++	C++	Src/Bin	Free	-	X	-	X	-	X	-	X	-	-	Library	
Faunus	Aurelius	Apache	x86	Java	Src	Free + Support	Both	-	X	X	X	-	Gremlin	X	X	-	Hadoop	
Gephi	Gephi Consortium	GPL/CDDL	x86	Java,OpenGL	Src/Bin	Free	-	X	-	X	-	X	-	X	X	-	Toolkit	
Giraph	Apache	Apache	x86	Java	Src	Free	-	X	-	X	X	X	-	X	-	-	BSP	
GraphStream	University Le Havre	LGPL/CeCILL-C	x86	Java	Src/Bin	Free	-	X	-	X	-	X	-	X	-	-	Library	
Hama	Apache	Apache	x86	Java	Src	Free	-	X	-	X	X	X	-	X	-	-	BSP	
MTGL	Sandia NL	BSD	XMT	C++	Src	Free	-	X	-	X	-	X	-	X	-	-	Library	
NetworkX	Los Alamos NL	BSD	x86	Python	Src/Bin	Free	-	X	-	X	-	X	-	X	-	-	Library	
PEGASUS	CMU	Apache	x86	Java	Src/Bin	Free	-	-	X	X	X	X	-	-	X	-	Hadoop	
STINGER	GT / GTRI	BSD	x86*/XMT	C	Src	Free	-	X	-	X	-	X	-	X	X	X	Library	
uRika	Yarc Data / Cray	Proprietary	XMT	Likely C++	Bin	\$\$\$\$	?	X	-	X	-	-	SPARQL	-	X	X	Appliance	

\* see page 2 of the paper


\* DEX is now Sparksee



# Evaluation


- **Four** fundamental graph kernels
  - Prefer a particular implementation
  - **Emphasize common** programming / traversal **styles**
  - Measure performance in multiple ways
    - **Time** to completion
    - **Memory** in use (MemoryMXBean for Java, OS reporting otherwise)
    - Qualitative analysis of **capabilities** and development **experience**
  - If not provided by package, written by **naïve programmer**
- Same R-MAT graphs used with all packages
  - Vertices: 1K (tiny), 32K (small), 1M (medium), and 16M (large)
  - Edges: 8K (tiny), 256K (small), 8M (medium), 128M (large)
  - Graphs of 1B edges considered, but not included due to number of systems that couldn't work with large






# Evaluation

- Single Source Shortest Paths (SSSP)
  - Must be a breadth-first traversal of the graph
  - Level-synchronous parallel where possible
  - Output unweighted distances from source to all vertices
- Application
  - Used for routing and connectivity
  - Building block for other algorithms (e.g. betweenness centrality)
  - Graph 500 benchmark




# Evaluation

- Connected Components
  - Based on Shiloach-Vishkin
  - Edge-parallel label-pushing style
  - Global graph metric that touches all edges in the graph
- Properties
  - Not theoretically work efficient, but edge-parallel pattern is representative of a broad class of graph algorithms
  - Read-heavy with sparse writes
  - Cache friendly graph traversal with cache-unfriendly random label read / assignment








# Evaluation


- PageRank centrality algorithm
  - Vertex-parallel Bulk Synchronous Parallel (BSP) power-iteration style
  - Measure of influence over information flow and importance in the network
  - Representative of a broad class of iterative graph algorithms
- Update benchmark
  - Perform edge insertions and deletions in parallel
  - Random access and modification of the structure
  - Real-world networks are in constant motion as new edges and vertices enter the graph





# DISCLAIMERS

- We are **not expert** Java programmers, Hadoop tuners/magicians, Python gurus (maybe a few), Scala wizards
- We are some of the primary maintainers of the STINGER graph package
- The **code** for these implementations is **available** at [github.com/robmccoll/graphdb-testing](https://github.com/robmccoll/graphdb-testing)
  - Test code is **BSD**, package licenses vary
  - We apologize if you feel that your particular work has been slighted or misrepresented
  - If any of these packages belong to you, we **invite** and **encourage** you to improve on our results and submit them back to the repository




# Maximum Size Completed

Package	Tiny	Small	Medium	Large
boost	X	X	X	X
dex	X	X		
mtgl	X	X	X	X
neo4j	X	X		
networkx	X	X	X	
orientdb	X	X	X	
sqlite	X	X		
stinger	X	X	X	X
titan	X	X	X	
bagel	X	X	X	
pegasus	X	X	X	X
giraph	X	X	X	X


Determined by crashing, running out of memory, or running an algorithm for longer than 24 hours.


Exception: Dex had a 1 million element ( $|V| + |E|$ ) limit.



# Memory Usage

3 orders of magnitude difference in memory usage





# Page Rank Performance

## Small


Range 10,000x

## Medium (x32)

Range 10,000x  
40~90x slower  
Except Pegasus,  
Giraph






## Large (x1024)

Range 100x  
~20x slower  
Except Pegasus




# SSSP and Components

## Small vs. Medium









# Observations on Quantitative

- Big performance gaps, possible groups
  - High Performance (ms response)
  - Medium Performance (sub minute response)
  - Storage Performance
    - Perhaps suited for storage, but not analysis
    - Not language or system, just not written with graph algorithms in mind
- Some parallel, but many not
  - Graph libraries generally parallel
  - Not much parallelism in (single node) disk-backed DBs
    - Parallelism only increases random read and write
    - Bottleneck is disk bandwidth and latency
- Data scaling is hard to predict
- Full results available at:
  - [github.com/robmccoll/graphdb-testing](https://github.com/robmccoll/graphdb-testing)
  - [arxiv.org/abs/1309.2675](https://arxiv.org/abs/1309.2675)



# Qualitative Observation

- Multitude of input file formats supported
  - No universal support
  - XML, JSON, CSV, or other proprietary binary and plain-text formats
  - Trade-off between size, descriptiveness, and flexibility
  - Challenge for data exchange
  - Edge-based, delimited, and self-describing easily parsed in parallel, translated to other formats
- Same data, same algorithms, same HW, DB / library is the difference
- No consensus on ACID / Transactions
  - No clear result on its affect on speed
  - Not clear that the applications exist
- No consensus on query languages, models



# Moving Forward

- To advance HPC outside of our community, software must be accessible
- Some applications don't compile out of the box
  - Self-contained packages are easiest
  - Package management systems don't work well on nodes inside a cluster
  - Build problems drive users away
- Lack of consistent approaches to documentation
  - Well-written tutorials better than sparse full API doc
  - Lack of adequate usage examples
    - Inserting vertices and edges, querying for neighbors is not enough
    - What is the best way to traverse the structure?
    - How should data ingest, analysis, and extracting results be performed?
  - Highlights
    - [NetworkX](#) is very well documented, both examples and API doc
    - [Titan's API](#) doc is very complete, but **lacks** examples of real-world **usage**
    - [Bagel](#), [STINGER](#) **lack formal documentation** or have **fragmented** and occasionally inaccurate documentation

# Acknowledgment of Support

